

Shivam Garg

San Francisco, CA | shivamg2604@gmail.com | linkedin.com/in/shivammgarg | github.com/22shivam | shivamg.com | 8127786998

EDUCATION

Indiana University Bloomington

B.S. in Computer Science and Mathematics | GPA: 3.9

May 2025

- **Achievements:** Herman B. Wells Scholar (**full-ride scholarship**, ~12 awarded yearly), Ex-National Cricket Team Player (Singapore)
- **Relevant Coursework:** Machine Learning, Computer Vision, Distributed Systems, Algorithms, Operating Systems, Real Analysis, Abstract Algebra

EXPERIENCE

Software Engineer at AGI Labs

July 2025 – Present

- Built an **automated payments and invoicing system** via the Stripe API, **saving 15–20 hrs/week** in manual work for our client, MediaTeam, a marketing agency.
- Developed a scalable **ad-tech analytics pipeline** pulling spend/ROI data from Google Ads, Meta, and TikTok APIs; providing MediaTeam with **dashboards with <1% sync failure and 5-min data freshness**.

Software Engineer and Co-founder at Spindl

November 2021 – May 2022

- Launched **multi-chain crypto checkout platform** with NFT showcase features, enabling simple link-based transfers and embeddable merchant widgets. (React, Node.js, MongoDB, AWS)
- Deployed with **EC2 Auto Scaling + Docker**, CI/CD pipelines, and monitoring via **Prometheus/Grafana**.
- Drove adoption to **10k+ users** by pairing technical execution with early viral growth strategies.

Software Engineer and Co-founder at Streamling

July 2021 – November 2021

- Built a **decentralized video streaming platform** with creator tokenization using DeSo blockchain + Media Network; reduced egress costs by leveraging peer-to-peer routing.
- Secured ~\$6k in funding and acceptance into IU's Startup Accelerator.

PROJECTS

LLM-Powered Automated Podcast Clipping

- Built a full-stack platform for **auto-generating 10-15 short viral clips from long podcasts**, leveraging Whisper transcription + Gemini LLM for highlight detection and computer vision (active speaker detection) for speaker-aware framing.
- Designed a cloud-native processing pipeline with AWS S3, Inngest queues, and serverless GPUs on Modal, enabling asynchronous, **fault-tolerant video rendering with scalable GPU acceleration**.
- Implemented Stripe integration for usage-based billing and credit deduction system, handling payments, webhooks, and role-based quotas.

Reentrant Distributed Lock Service

- Designed a **fault-tolerant distributed lock manager** in Python on top of an etcd cluster, supporting reentrant locks, auto-renewal, and structured logging.
- Ensured safe concurrency in microservices and job scheduling with a 3-node deployment **benchmarked at 5k+ ops/sec and <40ms latency**.
- Delivered with Docker Compose setup and 50+ test suite, validating lease expiry, contention, and multi-client exclusivity.

TECHNICAL SKILLS

- **Languages:** Python, C, Javascript, TypeScript, SQL, Java, HTML/CSS
- **Frameworks:** React, React Native, Node.js, Next.js, PyTorch, OpenCV, Jest, PyTest
- **Cloud & Tools:** AWS, Azure, Docker, Kubernetes, Git, PostgreSQL, MongoDB